

Research on Improvement of Dynamic Load Balancing in MongoDB

Xiaolin Wang, Haopeng Chen, Zhenhua Wang

School of Software

Shanghai Jiao Tong University

Shanghai, P.R.China

qiuye454@163.com, chen-hp@sjtu.edu.cn, aspiration@foxmail.com

Abstract—As a representative of NO-SQL database, MongoDB is widely preferred for its automatic load-balancing to some extent, which including distributing read load to secondary node to reduce the load of primary one and auto-sharding to reduce the load on specific node through automatically split data and migrate some of them to other nodes. However, on one hand, this process is storage-load based, which can't meet the demand due to the fact that some particular data are accessed much more frequently than others and the "heat" is not constant as time going on, thus the load on a node keeps changing even if with unchanged data. On the other hand, data migration will bring out too much cost to affect performance of system. In this paper, we will focus on the mechanism of automatic load balancing of MongoDB and propose a heat-based dynamic load balancing mechanism with much less cost.

Keywords—MongoDB; heat-based; load balancing; resource management; auto-sharding; cloud computing

I. INTRODUCTION

Recent years have witnessed a big explosion of unstructured and semi-structured data that are generated by the rapid development of internet. To store these large amount of data efficiently, Cloud computing, as a current commercial offering, started to become apparent in late 2007 [1]. Including Amazon S3 [2], Google Cloud Storage [3], Microsoft Azure [4] and so on. In this context, NO-SQL storage including Google's Bigtable [5] and its open-source implementation HBase [6], Amazon's Dynamo [7], Facebook's Cassandra [8], and LinkedIn's Voldemort [9] are becoming the focus of attention and large-scale distributed system is chosen to meet the need of storage capacity. In large-scale distributed system, fast response and high availability are two important performance objectives pursued by end users and applications. These performances are largely determined by the data allocation strategies and the load condition in the system because evenly distributed workload can help to optimize resource utilization, maximize the throughput and eliminate the potential overload. However, on the one hand, most load balancing for NO-SQL is not that mature or with low adaptability because of its short time development, on the other hand, the existing load balancing strategies for relational database are not fit NO-SQL database, thus we should develop new load balancing strategy which is more efficient and with higher applicability.

To design a proper load balancing strategy, three points should be taken into consideration. Firstly, due to the fact that some particular data are accessed much more frequently than others and the "heat-diffusion" is not constant as time going on, thus the load on a node keeps changing even if with unchanged data. So the load balancing strategy should be

heat-based, the heat in load balancing is the frequency of processing of the data. Secondly, load balancing should be a life long journey rather than just allocating the data when adding data to database, in other words, the location of data should not be immutable. Thirdly, when come to data migration, the amount of data together with migrate distance and other factors should be taken into account to decide the cost of migration. Besides, replica set should be used to share load except to ensure automated failover.

Automatic load balancing is an ideal target of database design, many storage products try to obtain it. However, they are mostly storage-based. For example, Dynamo achieves automatic load balancing mainly through "uniform hashing". The solution of load balancing in Bigtable is based on traditional server-farm. That is a master server monitors the load conditions of TServers continuously, and then migrates data from the overload server to an underloaded one. MongoDB [10] is precisely widely favored for the reason that it can achieve read and write extension automatically. Read extension is achieved mainly by replica set. Write extension is achieved mainly by auto-sharding. Sharding refers to the process of splitting data up and storing different portions of the data on different machines. It migrates chunks among different shards; each represents a smaller range of values within the shard's range. However, in our opinion, the load balancing strategy in MongoDB has at least three aspects that are not perfect. Firstly, MongoDB determines imbalance based solely on the difference of chunk number of each shard, in other words, it identifies imbalance only based on the amount of data. Thus it does not meet the common scene of the uneven distribution of hot spots. Secondly, MongoDB conducts chunk migration only taking the amount of data into consideration, without considering the physical distance and other specific shard's load information. Thirdly, MongoDB only has the scale up mechanism while scale down has not been involved in, so it needs to be optimized.

In this paper, we will take MongoDB for example to illustrate a new algorithm aiming to achieve automatic sharding based on heat diffusion and make full use of replica set to share load and ensure high availability at the same time. Section 2 reviews some related work and discusses the motivation of this work. Section 3 addresses the overview of the framework of our load balancing strategy. Then in Section 4, we propose the load balance strategy with detailed description of the algorithm. Experimental evaluation is presented in Section 5 with some simulation results. Finally Section 6 concludes the paper. Section 7 will give an acknowledgement to illustrate some additional information.

II. RELATED WORK AND MOTIVATION

Some of previous works have focused on load measurement and finding sources of imbalance. Efficient, scalable measurement of load [11] identifies whether load imbalance is a problem for a particular application. If the imbalance is a problem, then try to fix it, otherwise, ignore it. In [12], they try to find the cause of imbalance. Imbalance attribution [13] provides insight into the source code locations that cause of imbalance. However, these two works need a good understanding of application elements, which is too hard to achieve in most scenario. In [14], they introduce a dynamic and integrated resource scheduling algorithm (DAIRS), which treats CPU, memory and network bandwidth integrated for both physical machines and virtual machines and develop integrated measurement of the total imbalance level of a Cloud datacenter as well as the average imbalance level of each server for performance evaluation. In [15], they model the VE by dividing it regularly into a large number of square cells, each cell contains some objects, and its load is determined by the number of objects inside it, the author in [16] use virtual IDs to refer to the performance and capacity of nodes. Different IDs can handle different load levels and different processing time helps to find the idle servers, but these are storage-based measurement, which is not appropriate unless with even access.

Other researchers have focused on the data migration problem in a load rebalancing scenario. They try to identify a better method according to the cost or the speed of data migration. For example, in [17], two interference-aware prediction models are built to predict the migration time and performance impact for each action using statistical machine learning and then create a cost model to strike a right balance between the two ingredients of cost; H. C. Lim et al. [18] try to address the problem of automatic control for elastic storage. Two controllers are responsible for adding (or removing) a storage node and controlling the data migration respectively. However, these two works only pay too much attention on network resources, without considering which part of data to migrate and its impact on the operation on the database. In [19], they propose a hybrid control strategy for load balancing. This strategy would use the centralized load balancing strategy to redistribute the load of the node. From the global point of view, a controllable node controls two storage node clusters, and each storage node cluster belongs to two controllable nodes. Load balancing between the overlapping nodes continues timely operating the iterative batch load method to distribute the load to the global. But it is not a real-time method.

Actually, there are other researches on the improvement of rebalancing on MongoDB. For example, in [20], for each chunk in MongoDB, its load determined by the number of operations including adding, deleting and querying on it. However, the weights of different operations are hard to assign only by experience.

In conclusion, heat-based load balancing of NO-SQL database is necessary for large-scale distributed system. However, there is not a comprehensive solution has been proposed, especially a method combined with architecture of database, like using replica set to reduce the rebalancing cost. Thus a better method

should be proposed.

III. FRAMEWORK

We reformed the framework of MongoDB, especially the decision making group as follow, though it is based on the framework of MongoDB, it is widely applicable.

Our load balancing strategy can balance the workload among different physical servers and virtual servers based on heat diffusion in a dynamic way. Fig. 1 features the core components in our load balancing strategy as follows and mainly reformed decision making group and added load agent to replica management. Thus we will describe the first two parts briefly and the third part in detail.

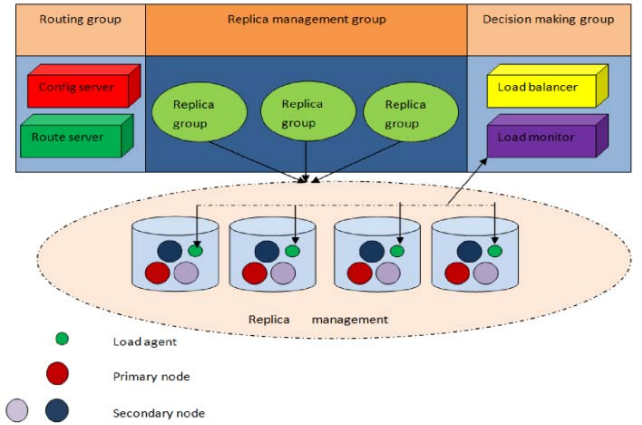


Fig. 1. Framework

A. Description of Routing Group

As we all know, in most NO-SQL database including MongoDB, different parts of data are partitioned to different nodes, thus, each time when there is a client request, it is routed to routing group to decide which node to get the specific data. That is, routing group act as a proxy between client and storage system. Besides, in MongoDB, if it is only a read request, routing group will map it to a secondary node to share the load of primary one.

Config server: It does not only keep the map between key range and VM server, but also keep a map between VM server and physical server.

Route server: For updating or inserting, route server will route request based on key and the information of config server, for querying, it will route request to a secondary node if the criteria contains specific key. Otherwise, it does scatter and gather on secondary node.

B. Description of Replica Management Group

Data Nodes: As described in figure 1, on one hand, the data are distributed on multiple shards; each shard acts as a replica group, including one primary node and one or several second nodes. Primary node is in charge of writing request, updating request, and read request, but only read request can act on secondary node. On the other hand, chunks are hosted in shards, each chunk specifies a key range, and chunks together span the whole key space.

Load Agent: The purpose of load agent is to provide load information of the server, physical or virtual; it provides information to the load monitor. Because different replica group might have replicas in the same physical server, it is hard to presume the percentage of resource is being used by which part at particular moment. Therefore, we think it is better to allocate one replica on each VM. In addition, it is easy to operate and the cost is reduced when come to data migration.

C. Description of Decision Making Group

The Decision making group acts as an important part in loadbalancing system and keeps load evenly across the cluster. It is in charge of making decision whether to take load balancing steps according to our load balancing policy. In MongoDB, loadbalancing strategy is to transfer data in chunk unit across shards according to the number of chunks, but in our reformed framework, we try to avoid data migration which brings lots of cost and with the help of load monitor, rebalancing decision is based on the load on each node.

Load Monitor: Load monitor is a new part in the framework; it collects load information from every load agent within certain time interval. The load information should be refreshed at a suitable interval so that the information provided is not expired. In addition, it will compare the load with the normal load range, if the load is found to be abnormal, and then it will pass a sign to load balancer.

Load balancer: load balancer is a reconstructed part by us; it is in charge of taking rebalancing steps according to rebalancing workflow and algorithms. Whenever load balancer is aroused by Load Monitor, firstly it will figure out whether the node is a physical node and try to rebalance the cluster according to specific workflow. In MongoDB, the migration chunk is just the top chunk of each shard, ignoring the heat on it. But in our framework, the chunk with largest heat will be split to avoid failing to distribute the access load. Thus; it works as activist in a core part.

IV. HEAT DIFFUSION BASED LOAD BALANCING WORK FLOW AND ALGORITHMS

Routing group in MongoDB could be viewed as a normal shard, because it is also constituted with a replica set, and the load on it will be desperate to one of idle node or other routing node will less load. Thus, we will describe how to define overloaded and underloaded node with specific algorithm at first, then we will describe the improved automatic sharding from two parts, VM overloaded balancing and physical overloaded balancing and illustrate the procedure with flow diagrams, and then underloaded process will be described finally.

A. Exception Detection Algorithm

The overload criterion is so important that we take many experiment to define it. We used the utilization of each node to evaluate the load on it. Actually, there are many types of resources on either physical node or virtual node, such as CPU, memory, bandwidth, disk and so on. However, we choose the first three resources to define the load on each node for simplification. To describe the algorithms, we will define all parameters as follows.

TABLE I
PARAMETERS OF LOADBALANCING

| Parameters | Description |
|------------|--|
| U_{cni} | The monitored utilization of CPU of node i |
| U_{mni} | The monitored utilization of memory of node i |
| U_{bni} | The monitored utilization of bandwidth of node i |
| U_{cu} | The upper bound of utilization of CPU |
| U_{mu} | The upper bound of utilization of memory |
| U_{bu} | The upper bound of utilization of bandwidth |
| U_{cl} | The lower bound of utilization of CPU |
| U_{ml} | The lower bound of utilization of memory |
| U_{bl} | The lower bound of utilization of bandwidth |
| W_c | The weight of utilization of CPU |
| W_m | The weight of utilization of memory |
| W_b | The weight of utilization of bandwidth |
| U_{ii} | The integrated utilization of node i |
| U_l | The lower bound of utilization of node i |

Firstly, we defined the upper bound of indexes, U_{cu} , U_{mu} , U_{bu} , and lower bound of three resources of each node. Then we define overloaded and underloaded node by comparing the monitored utilization with its upper bound and lower bound.

Because when the utilization of any resources on the node is over its upper bound, the availability on it will be seriously affected, but to identify whether a node is underloaded, we should take monitored utilization of each resource into consideration to get an integrated measurement.

$$U_{ii} = W_c * U_{cni} + W_m * U_{mni} + W_b * U_{bni} \quad (1)$$

$$U_l = W_c * U_{cl} + W_m * U_{ml} + W_b * U_{bl} \quad (2)$$

B. VM Overloaded Balancing

As described in (a), when specific VM is detected overloaded, it means that there is at least one hot point on the node. Thus firstly we locate specific hot chunk according to the number of request on it, then split the chunk at the middle of the key range of the chunk, this step will avoid failing to distribute the access load and initiate the auto-sharding of MongoDB, as discussed in [10], whenever the difference of chunk number is over 2, which could be set, the auto-sharding will be started. If the condition is not resolved yet, the process will be done one more time until the load on all VM nodes turned to be normal.

C. VM underloaded Balancing

For underloaded VM node, we wouldn't process it at once because it is possible that an overload one will find this node to merge after a while. So we will start a timer and wait for a specific interval. If until time up, there is no request, we could judge there is no proper overload node in the cluster. Then we will try to find a paired underloaded node. If succeed, we will migrate the chunks on the shard with less data to the paired underloaded VM, and after the migration operation, all of the indexes of both the underloaded VM node are not over the upper bound.

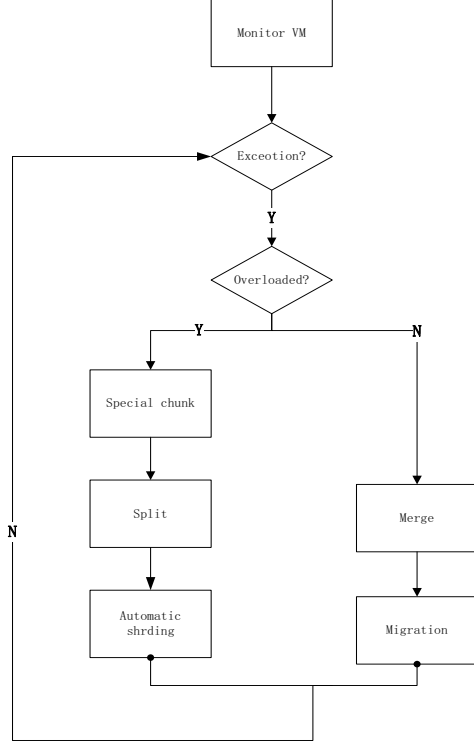


Fig. 2. VM load balancing work flow

D. Physical Overloaded Balancing

As described in (b), when specific physical node overloaded, we will not migrate data immediately, instead, we will try to change primary node to secondary one on the overloaded physical node which will reduce the cost of rebalancing substantially, because the write requests could be only executed on primary node, it is possible that the switch action could reduce the load on the physical load and lead it to a normal one.

TABLE 2
Parameters of Algorithm

| Parameters | Description |
|------------|---|
| L_{pi} | The load on primary node of replica set i |
| L_{ii} | The load on secondary node i of replica set i |
| R_i | Replica set of i |
| V_i | Virtual node i |
| P_i | Physical node i |
| L_i | Load on Physical node p |

TABLE 3
Algorithm

| |
|--|
| Algorithm: Algorithm |
| Inputs : L_{pi} on V_p of P_i , L_{ii} on V_i , R_i of P_j |
| Output : result:enum<1, 0> |
| For V_i in $R_i(i=1,2,3,...)$ |
| if L_i is lowest in R_i |
| $a = L_i - L_{pi} + L_{ii}$ |
| $b = L_j + L_{pi} - L_{ii}$ |
| a is in normal interval |
| and b is in normal interval |
| then result = 1 |
| else |
| result = 0 |
| return result |

Firstly, we will predict whether the switch action could resolve the problem according to predict algorithm above, if it can achieve the target, then we switch the primary node, that is, the primary node is switched to second node of the replica set, while the secondary node is switched to primary one. If it isn't succeed, we should try to search a paired physical node with more spare resource and migrate a specific VM to it, the data on the VM must be as small as possible. After the migration operation, either the overloaded physical node or the paired one is normal. However, if there is no physical node could be added, we should add a new physical node to the cluster, and specific VM will be migrated to it by paired operation. Actually, in most cases, we could just modify the map between physical node and VM node without data migration.

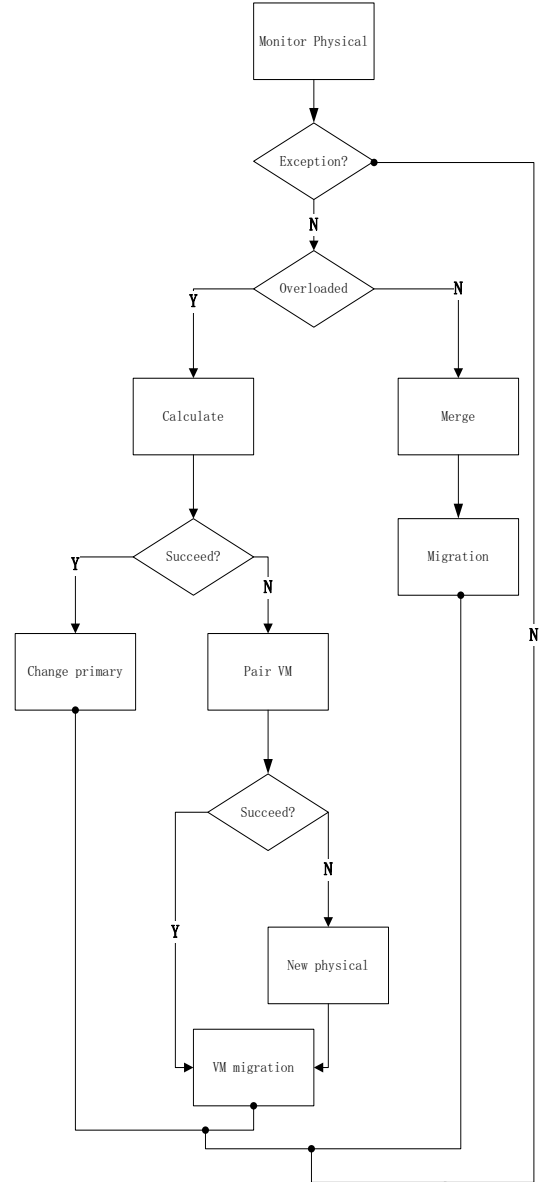


Fig. 3. Physical load balancing work flow

E. Physical Underloaded Balancing

For underloaded physical node, we wouldn't process it at once because it is possible that an overload one will find this node to merge after a while. So we will start a timer and wait for a specific interval. If until time up, there is no request, we could judge there is no proper overload node in the cluster.

For physical node, we will search for other underloaded node to pair, if succeed to find one, and then migrate the node with less data to another one to merge them. Thus, one of them will become idle. Otherwise, restart the timer and wait for another pair request.

V. EXPERIMENTS

A. Experiment Configurations

To demonstrate the effectiveness of the propose framework, we build an experiment environment with 8 physical nodes and 25 virtual nodes.

In the physical nodes, XenServer is set up as the virtualization server. And in each virtual node, our load balancing algorithm is deployed and interacts with the API of XenServer.

In the virtual nodes, Ubuntu 12.04 Server is installed in each node. And shards as MongoDB storage nodes are deployed in these virtual nodes.

The details of the physical machines used to set up the experiment environment are as follows.

Table 5.The detail of physical machine

| Category | Name | CPU | Memory (GB) | Disk (GB) |
|-----------------------|------------|-------------------------|-------------|-----------|
| Virtualization Server | XenServer1 | Intel i5 3.30G Hz | 4 | 500 |
| | XenServer2 | | 4 | 500 |
| | XenServer3 | | 4 | 500 |
| | XenServer4 | | 4 | 500 |
| Proxy Node | Proxy1 | Intel i3 3.30G Hz | 4 | 500 |
| | Proxy2 | | 4 | 500 |
| Client | Client1 | | 4 | 500 |
| | Client2 | | 4 | 500 |

The details of the values of the parameters mentioned above for work-load analysis are as follows.

Table 4.Configuration of parameters for work-load analysis

| Parameters | Values for physical nodes | Values for virtual nodes |
|--------------------------|---------------------------|--------------------------|
| opt of CPU | 0.6 | 0.7 |
| opt of Memory | 0.95 | 0.95 |
| opt of NetworkIO | 0.95 | - |
| weight of CPU | 0.1 | 0.2 |
| weight of Memory | 0.2 | 0.8 |
| weight of NetworkIO | 0.2 | - |
| threshold of underloaded | 0.3 | 0.3 |

The choice of the parameters should weight the costs and performance. Presently, we choose the parameters according to

the performance priority principle to guarantee the storage system performance. And the method for determining the values of parameters will be presented by the later work.

Physical machine Clinet1 and Client2 are used to simulate clients for data access. 200 users are simulated by them respectively. At the beginning of the experiment, there is no user access to the storage system. It lasts for 1000s and the work-load of the entire system is quite light in the first 1000s. Then, the simulation of user access starts. The number of concurrent simulated users increases evenly in the first 1000s. Then it reaches steady state. The simulated users send their data requests continuously and the interval between two requests is 100ms. The state lasts for 5000s then the simulation stops. we use ycsb as test tools, its an effective NoSQL benchmark tool.

B. Resource Management Impact Evaluations

Ideally, when the system is busy, more physical machines should be in service to guarantee the basic performance of the entire system. However, when the system is free, some free physical machines should be shut down or go to sleep for energy-saving. As the follow figure 4 shows, at the beginning of the experiment, we deliberately set the number of physical nodes in service to be 4. But the work-load of the entire system is extremely low since there is no simulated user access to the system at the beginning. The resource management algorithms work and the number of physical machines in service is reduced from 4 to 3 for energy-saving. Then the user access simulation starts. Thus the work-load of the system gets heavier and heavier. At this time, the resource management algorithms work and during the heavy work-load period, the number of physical nodes in service is increased from 3 to 4 to guarantee the basic performance of the system. Since the work-load caused by simulation user access is not steady, the number of physical nodes in service changes.

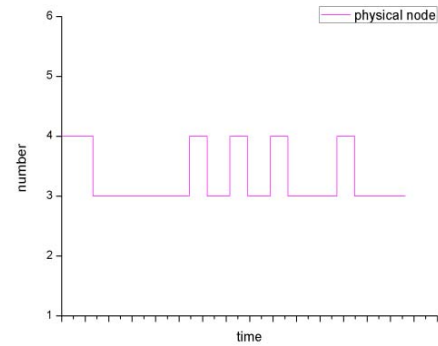


Fig. 4.Physical node in the cluster

C. Dynamic Load Balancing Impact Evaluations

When the work-load exception is detected, the dynamic load balancing algorithm begins to work. As figure shown, this physical node is extremely heavy during the work-load heavy period. Thus, its resource utilization exceeds the defined threshold. Fortunately, with the impacts of the dynamic load balancing algorithms, its extremely heavy work-load is reduced as shown in the figure.

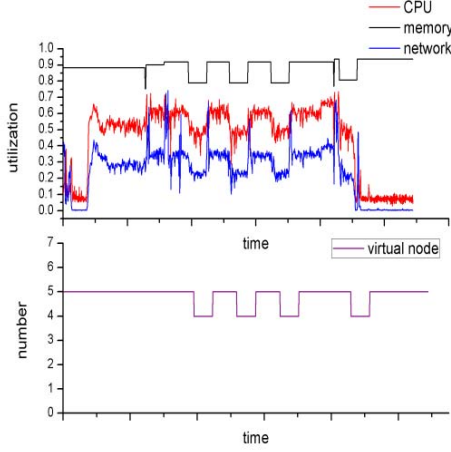


Fig. 5. Physical node load balancing

In addition, the work-load is transferred to another physical node. As shown in the follow figure6, at the beginning, this physical node is determined to be free and with the impact of the resource management algorithm, all of its virtual nodes have been migrated to other physical nodes. But because of the heavy work-load caused by the simulation user, some virtual nodes are moved to it to guarantee its basic performance. Thus its work-load is increased during the heavy work-load period.

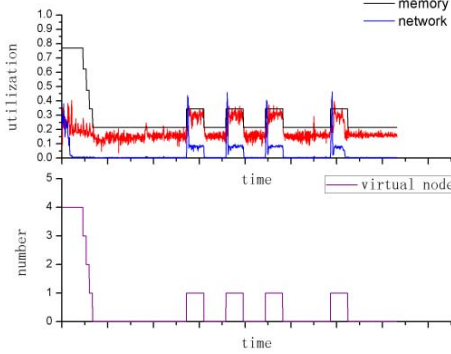


Fig. 6. xen-p2 resource utilization and the number of virtual nodes over time

And as the follow figures show, all the computation resource utilization is controlled under its defined threshold. It further demonstrates the effective of the dynamic load balancing algorithm. The benefit is that it can avoid the system bottlenecks caused by the excessive utilization of computation resource.

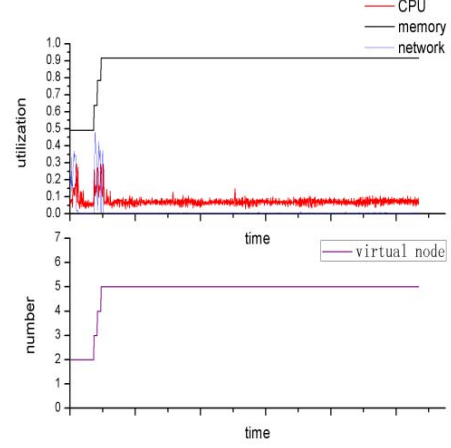


Fig. 7. xen-p3 resource utilization and the number of virtual nodes over time

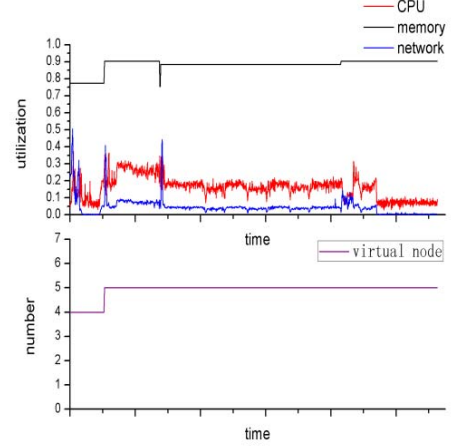


Fig. 8. xen-p4 resource utilization and the number of virtual nodes over time

VI. CONCLUSION

As discussed above, dynamic heat diffusion-based load balancing is important to storage application with hotspot data. Traditional data amount-based load balancing cannot meet the need in real scenario, including the load balancing mechanism of MongoDB. In addition, replica set is more than just a backup; it also can help to balance the load in the system.

At presented before, our work is effective to some extent. It effectively rebalanced the load in the system based on heat-diffuse, and made full use of replica set to reduce the migrate cost.

In addition, our mechanism is not only fit to MongoDB, it could be used into other database to optimize the performance of load balancing.

VII. ACKNOWLEDGEMENT

This paper is supported by Shanghai Municipal Science and Technology Commission under Grant No.11dz1502500.

REFERENCES

- [1] IBM. IBM Introduces Ready to Use Cloud Computing. IBM PressRelease, 15th November 2007.
- [2] Amazon, Amazon S3 [OL], available at: <http://aws.amazon.com/s3>.
- [3] Google, GoogleCloudStorage [OL], available at: <http://www.google.com/enterprise/cloud>.
- [4] Microsoft, Microsoft Azure [OL], available at: <http://www.windowsazure.com/en-us/>
- [5] Amazon SimpleDB [OL], available at: <http://aws.amazon.com/simpledb/>.
- [6] HBaseHomepage [OL], available at: <http://hbase.apache.org>.
- [7] G. DeCandia et al. Dynamo: Amazon's Highly Available Key-value Store. In SOSP, page 220, 2007.
- [8] A. Lakshman and P. Malik. Cassandra: A Decentralized Structured Storage System. SIGOPS Oper. Syst. Rev., 44(2):35 {40, 2010}.
- [9] Project Voldemort [OL], available at : <http://project-voldemort.com>.
- [10] MongoDB.features. [OL], available at : <http://www.mongodb.org/>
- [11] K. A. Huck and J. Labarta. Detailed load balance analysis of large scale parallel applications. In Par.Processing, 2010. pp.102-110
- [12] N. R. Tallent, J. M. Mellor-Crummey, M. Franco, R. Landrum, and L. Adhianto. Scalable fine-grained call path tracing. In Intl. Conf. on Supercomputing, 2011..pp. 123-125.
- [13] N. R. Tallent, L. Adhianto, and J. M. Mellor-Crummey. Scalable identification of load imbalance in parallel executions using call path profiles. In Supercomputing, 2010. pp. 103-112.
- [14] Wenhong Tian, Yong Zhao, University of Electronic Science and Technology of China, Chengdu, China A DYNAMIC AND INTEGRATED LOADBALANCING SCHEDULING ALGORITHM FOR CLOUD DATACENTERS Proceedings of IEEE CCIS2011. pp. 213-220.
- [15] Deng, Yunhua and Lau, Rynson W. H. (2010): Heat diffusion based dynamic load balancing for distributed virtual environments. In: Proceedings of the 2010 ACM Symposium on Virtual Reality Software and Technology 2010. pp. 203-210.
- [16] Lin Xia, Han-Cong Duan, Xu Zhou, Zhifeng Zhao, Xiao-Wen Nie, "Heterogeneity and Load Balance in Structured P2P Syste" 2010. pp.303-310.
- [17] Yonghui Zhang, Chunhong Zhang, Yang Ji, Wei Mi, "A Novel Load Balancing Scheme for DHT-BASED Server Farm", 2010 3rd IEEE International Conference on Broadband Network and Multimedia Technology (IC-BNMT), 2010, pp. 980-984.
- [18] Xiulei Qin, Wenbo Zhang, Wei Wang, Jun Wei, Xin Zhao, Tao Huang: Towards a Cost-Aware Data Migration Approach for Key-Value Stores. CLUSTER 2012: 551-556
- [19] H. C. Lim, S. Babu, J. S. Chase. Automated control for elastic storage. In: Proc. of the 7th international conference on Autonomic computing (ICAC '10). 2010. pp. 1-10.
- [20] Yimeng Liu ; Sch. of Comput. & Inf. Technol., Beijing Jiaotong Univ., Beijing, China ; Research on The Improvement of MongoDB Auto-Sharding in Cloud Environment. Computer Science & Education (ICCSE), 2012 7th International Conference on Date of Conference: 14-17 July 2012 . pp.851 - 854